

```
In [36]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df= pd.read_excel("DA.xlsx")
```

```
In [3]: df.head()
```

```
Out[3]:
```

	Sr no.	Random	Vehicle Model	Price	Manufacturing year	Distance driven (in miles)	Engine Type	Engine(in litres)	Power	No of owners
0	1	0.941111	BMW 5 Series	16300	2019	72000	Diesel	2	252 PS	1
1	2	0.139268	BMW 5 Series	17490	2019	50000	Diesel	2	252 PS	1
2	3	0.180856	BMW 5 Series	20990	2019	59500	Petrol	2	181 BHP	2
3	4	0.611808	BMW 5 Series	19250	2019	54500	Diesel	2	190 PS	1
4	5	0.126623	BMW 5 Series	13450	2019	46950	Petrol	2	188 BHP	2

```
In [4]: df.drop('Sr no.', axis=1, inplace= True)
```

```
In [5]: df.head()
```

```
Out[5]:
```

	Random	Vehicle Model	Price	Manufacturing year	Distance driven (in miles)	Engine Type	Engine(in litres)	Power	No of owners
0	0.941111	BMW 5 Series	16300	2019	72000	Diesel	2	252 PS	1
1	0.139268	BMW 5 Series	17490	2019	50000	Diesel	2	252 PS	1
2	0.180856	BMW 5 Series	20990	2019	59500	Petrol	2	181 BHP	2
3	0.611808	BMW 5 Series	19250	2019	54500	Diesel	2	190 PS	1
4	0.126623	BMW 5 Series	13450	2019	46950	Petrol	2	188 BHP	2

```
In [6]: df.drop('Random', axis=1, inplace= True)
```

```
In [7]: df.head()
```

Out[7]:

	Vehicle Model	Price	Manufacturing year	Distance driven (in miles)	Engine Type	Engine(in litres)	Power	No of owners
0	BMW 5 Series	16300	2019	72000	Diesel	2	252 PS	1
1	BMW 5 Series	17490	2019	50000	Diesel	2	252 PS	1
2	BMW 5 Series	20990	2019	59500	Petrol	2	181 BHP	2
3	BMW 5 Series	19250	2019	54500	Diesel	2	190 PS	1
4	BMW 5 Series	13450	2019	46950	Petrol	2	188 BHP	2

In [8]: `df.shape` *#For fetching total number of rows and columns*

Out[8]: (100, 8)

In [9]: `df['Power'].count()`

Out[9]: 100

In [10]: `df['Power'].value_counts()`

```
Out[10]: Power
181 BHP    15
252 PS     12
190 PS     12
265 PS      9
184 PS      8
248 BHP      6
188 BHP      5
394 PS      5
282 BHP      5
261 BHP      3
184 BHP      3
335 BHP      3
190 BHP      3
249 BHP      2
340 PS      2
252 BHP      1
264 BHP      1
187 BHP      1
192 BHP      1
389 BHP      1
265 BHP      1
369 BHP      1
Name: count, dtype: int64
```

```
In [11]: df['Engine Type'].value_counts()
```

```
Out[11]: Engine Type
Petrol    51
Diesel    49
Name: count, dtype: int64
```

```
In [12]: df['Engine_Power'] = df['Power'].str.extract('(\d+\.\d+|\d+)', expand=False).astype(float)
```

```
In [13]: df.head()
```

Out[13]:

	Vehicle Model	Price	Manufacturing year	Distance driven (in miles)	Engine Type	Engine(in litres)	Power	No of owners	Engine_Power
0	BMW 5 Series	16300	2019	72000	Diesel	2	252 PS	1	252.0
1	BMW 5 Series	17490	2019	50000	Diesel	2	252 PS	1	252.0
2	BMW 5 Series	20990	2019	59500	Petrol	2	181 BHP	2	181.0
3	BMW 5 Series	19250	2019	54500	Diesel	2	190 PS	1	190.0
4	BMW 5 Series	13450	2019	46950	Petrol	2	188 BHP	2	188.0

In [14]: `df.drop('Power', axis=1, inplace=True)`

In [15]: `df.head()`

Out[15]:

	Vehicle Model	Price	Manufacturing year	Distance driven (in miles)	Engine Type	Engine(in litres)	No of owners	Engine_Power
0	BMW 5 Series	16300	2019	72000	Diesel	2	1	252.0
1	BMW 5 Series	17490	2019	50000	Diesel	2	1	252.0
2	BMW 5 Series	20990	2019	59500	Petrol	2	2	181.0
3	BMW 5 Series	19250	2019	54500	Diesel	2	1	190.0
4	BMW 5 Series	13450	2019	46950	Petrol	2	2	188.0

In [16]: `df['Manufacturing year'].value_counts()`

Out[16]:

```

Manufacturing year
2019      38
2020      20
2021      20
2022      11
2023      11
Name: count, dtype: int64

```

In [17]: `df['Engine Type'].value_counts()`

Out[17]:

```

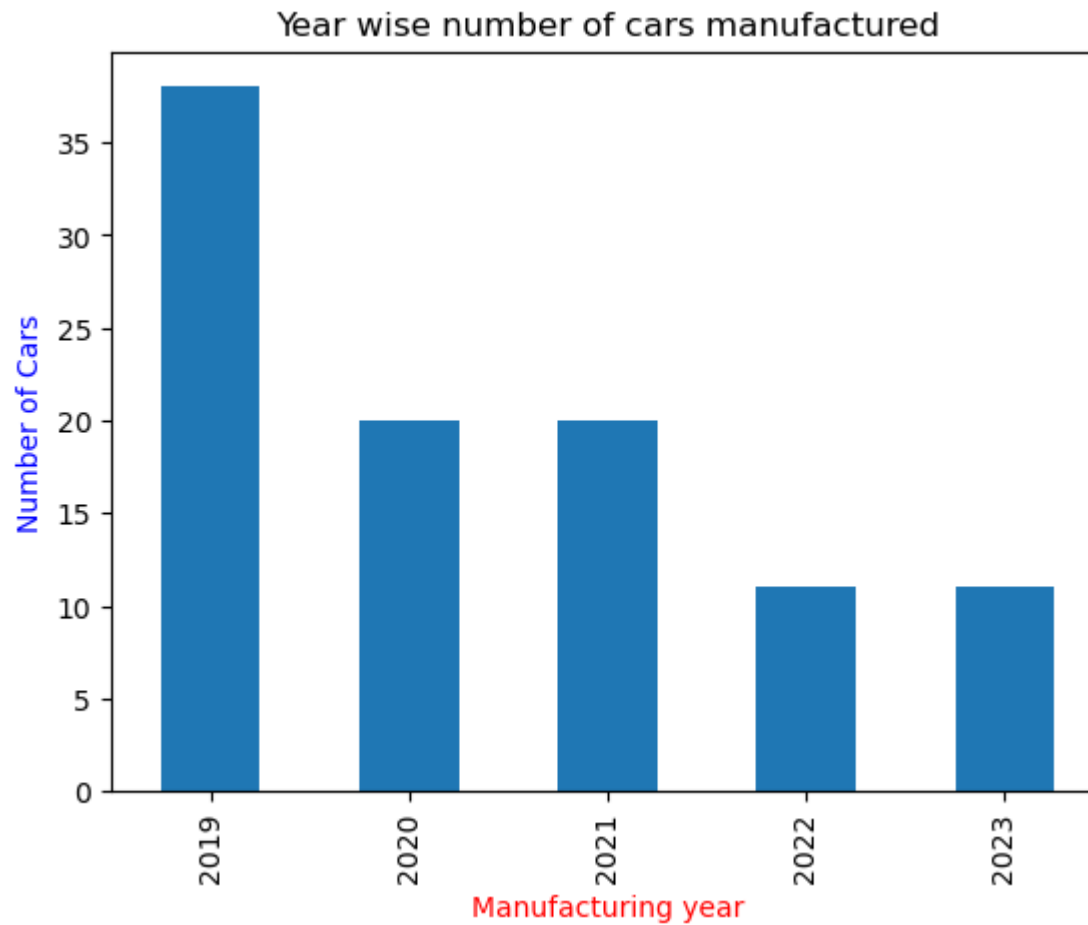
Engine Type
Petrol      51
Diesel      49
Name: count, dtype: int64

```

Q.2


```
In [19]: df['Manufacturing year'].value_counts().plot.bar()  
plt.xlabel('Manufacturing year', color= 'red')  
plt.ylabel('Number of Cars', color= 'blue')  
plt.title('Year wise number of cars manufactured')
```

Out[19]: Text(0.5, 1.0, 'Year wise number of cars manufactured')



```
In [21]: df.plot.scatter(x='Price', y='Engine_Power')
```

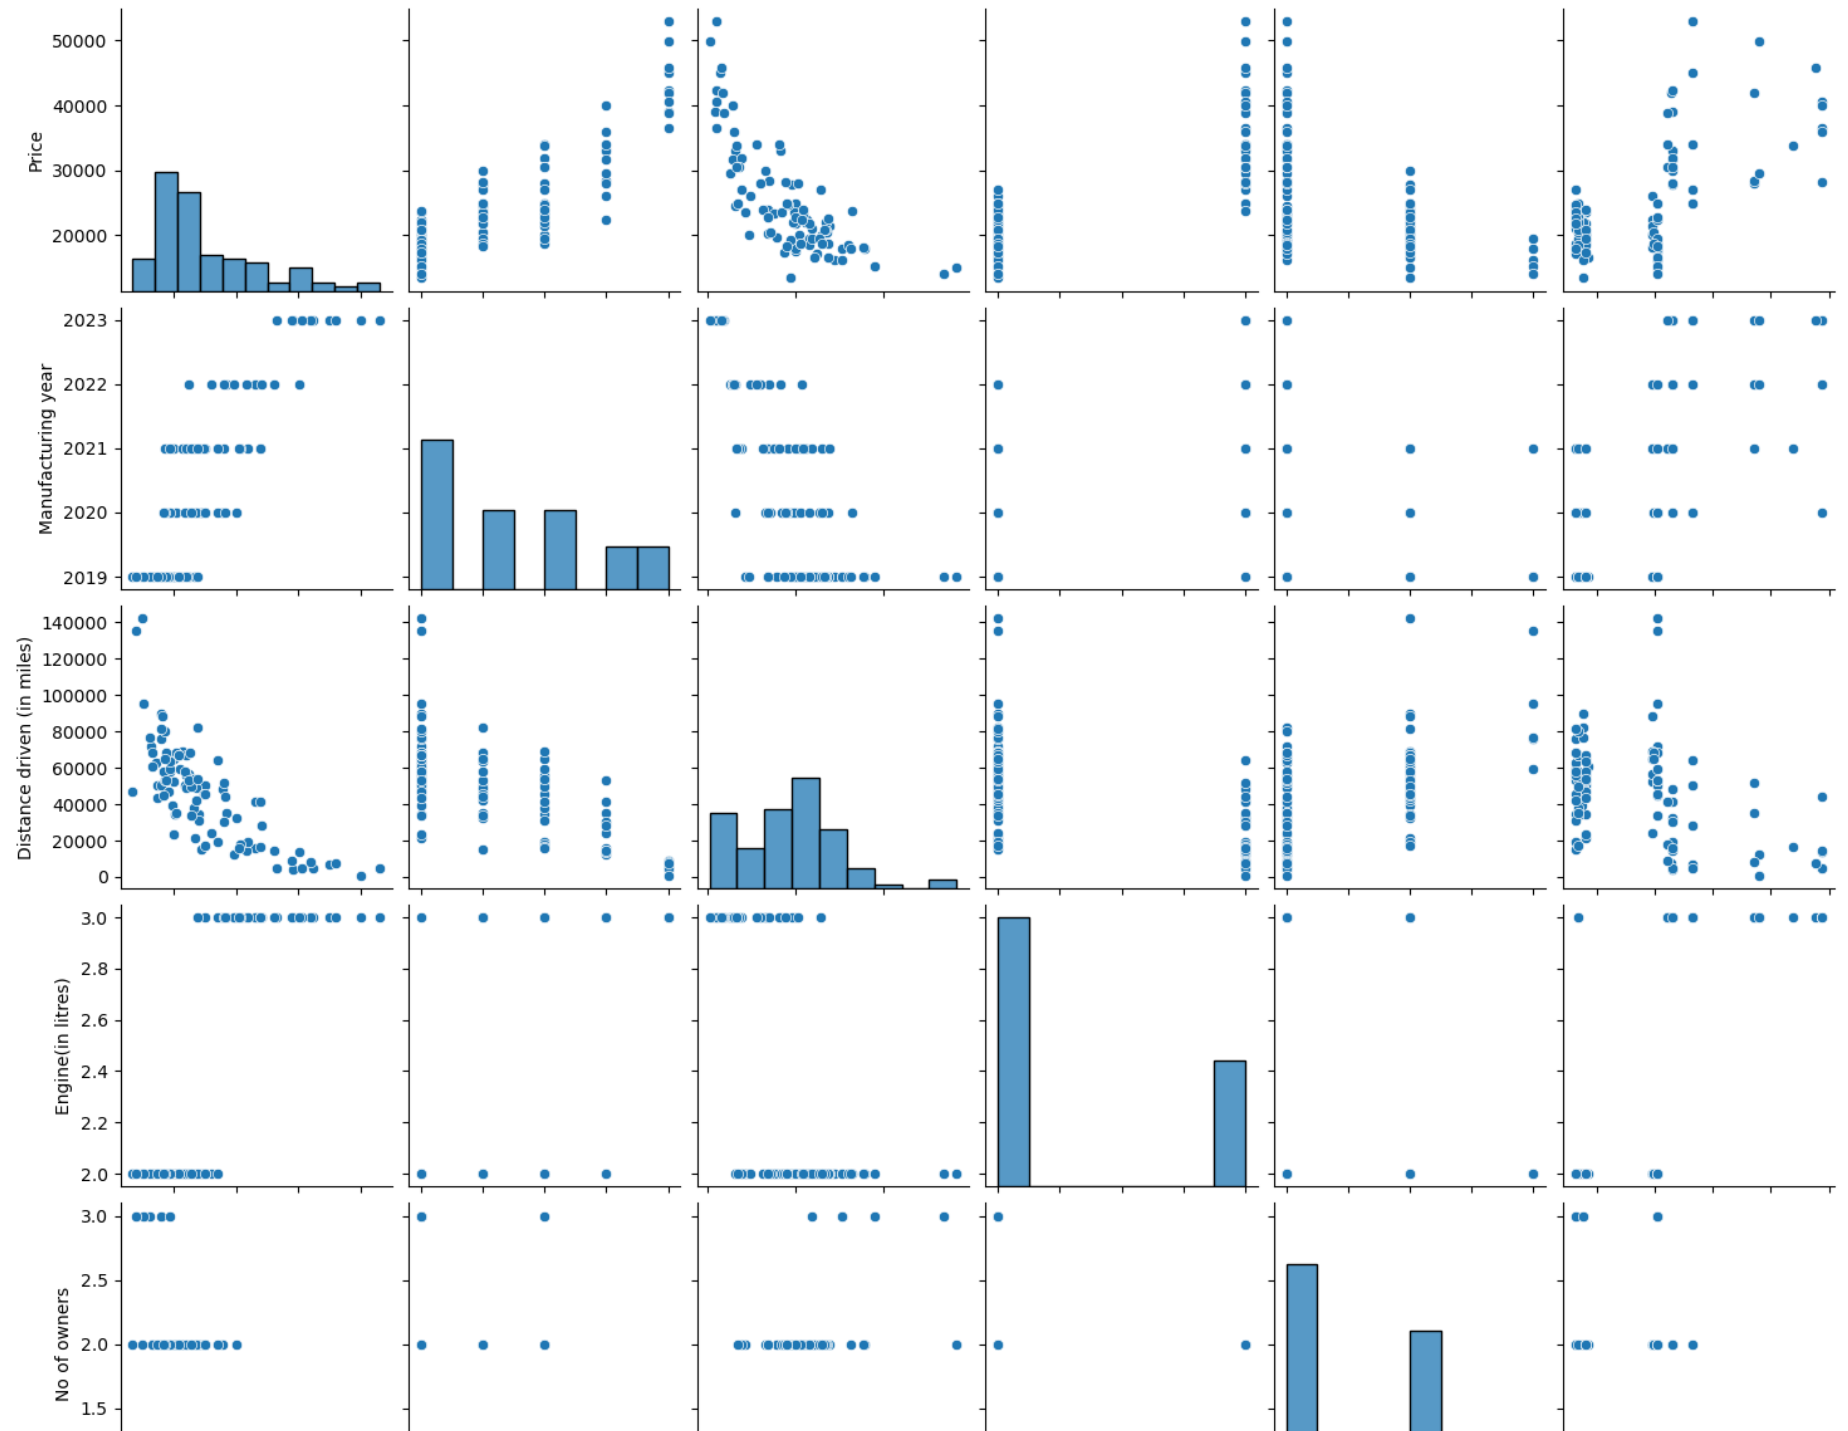
Out[21]: <Axes: xlabel='Price', ylabel='Engine_Power'>




In [22]: `sns.pairplot(df)`

C:\Users\apat\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning: The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)


Out[22]: <seaborn.axisgrid.PairGrid at 0x2b04f062910>






```
In [23]: df.plot.scatter(x='Manufacturing year', y= 'Distance driven (in miles)')
```

```
Out[23]: <Axes: xlabel='Manufacturing year', ylabel='Distance driven (in miles)'>
```

```
In [24]: sns.scatterplot(x=df['Distance driven (in miles)'], y=df['Price'], data=df)
```

```
Out[24]: <Axes: xlabel='Distance driven (in miles)', ylabel='Price'>
```



Q3.

```
In [27]: numeric_columns= ['Price', 'Distance driven (in miles)', 'Engine_Power', 'Manufacturing year']  
descriptive_stats= df[numeric_columns].describe().transpose()  
print(descriptive_stats)
```

	count	mean	std	min	25%	\
Price	100.0	25066.13	8328.729507	13450.0	19212.5	
Distance driven (in miles)	100.0	45632.53	26508.067132	904.0	24075.0	
Engine_Power	100.0	236.90	60.783354	181.0	184.0	
Manufacturing year	100.0	2020.37	1.375507	2019.0	2019.0	

	50%	75%	max
Price	22702.0	28261.75	52950.0
Distance driven (in miles)	48350.0	61361.50	142000.0
Engine_Power	248.0	265.00	394.0
Manufacturing year	2020.0	2021.00	2023.0

```
In [31]: mileage_bins= [0, 50000, 100000, float('inf')]
mileage_labels= ['0-50k', '50k-100k', '100k+']
df['Mileage Category']= pd.cut(df['Distance driven (in miles)'], bins=mileage_bins, labels= mileage_labels)
price_stats_by_mileage= df.groupby('Mileage Category')['Price'].describe()
print(price_stats_by_mileage)
```

	count	mean	std	min	25%	\
Mileage Category						
0-50k	56.0	29164.303571	8902.481039	13450.0	23168.75	
50k-100k	42.0	20105.166667	2751.272539	15277.0	18275.00	
100k+	2.0	14497.500000	703.571247	14000.0	14248.75	

	50%	75%	max
Mileage Category			
0-50k	27390.0	33965.00	52950.0
50k-100k	19495.0	21962.50	27990.0
100k+	14497.5	14746.25	14995.0

```
In [32]: df['Price'].describe()
```

```
Out[32]: count    100.000000
mean      25066.130000
std       8328.729507
min       13450.000000
25%      19212.500000
50%      22702.000000
75%      28261.750000
max       52950.000000
Name: Price, dtype: float64
```

```
In [33]: df['Price'].describe()
```

```
Out[33]: count      100.000000  
mean      25066.130000  
std       8328.729507  
min       13450.000000  
25%       19212.500000  
50%       22702.000000  
75%       28261.750000  
max       52950.000000  
Name: Price, dtype: float64
```

```
In [34]: df['Price'].std()
```

```
Out[34]: 8328.729507262888
```

Q4.

```
In [39]: mean_x=25066.13  
s=8328.72  
n=100  
confidence_level= 0.95  
  
#Z score for a 95% confidence interval  
  
z_score= 1.96  
  
#Calculate standard error  
standard_error= s/ np.sqrt(n)  
  
lower_bound= mean_x- z_score *standard_error  
upper_bound= mean_x + z_score *standard_error  
  
#create a dataframe to display the results  
confidence_intervals= pd.DataFrame({  
    'Confidence Level': [95],  
    'Lower Bound': [lower_bound],  
    'Upper Bound': [upper_bound]  
})  
  
#Display the result  
print(confidence_intervals)
```

	Confidence Level	Lower Bound	Upper Bound
0	95	23433.70088	26698.55912

Q5.

```
In [42]: import statsmodels.api as sm
import pandas as pd
from scipy.stats import ttest_1samp

X= df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power']]
X= sm.add_constant(X)

y=df['Price']

model= sm.OLS(y, X).fit()

predicted_prices= model.predict(X)

df['Predicted Prices']= predicted_prices

provided_averages= {'0-20000': 39494, '20000-40000': 36473, '40000-60000': 30222, '60000-80000': 36055, '80000-100000': 35027, '100000-120000': 30000}
for mileage_range, average_price in provided_averages.items():
    model_group = df[df['Distance driven (in miles)'].between(int(mileage_range.split('-')[0]), int(mileage_range.split('-')[1]))]
    t_statistic, p_value = ttest_1samp(model_group, average_price)
    print(f'Hypothesis Test for {mileage_range}, Mileage Range:')
    print(f'Test Statistic: {t_statistic}')
    print(f'P-value: {p_value}')

    if p_value < 0.05:
        print('Reject the null hypothesis: There is a significant difference.')
    else:
        print('Fail to reject the null hypothesis: There is no significant difference.')
    print('\n')
```

Hypothesis Test for 0-20000, Mileage Range:
Test Statistic: -24.302771688858385
P-value: 2.1906595415161028e-17
Reject the null hypothesis: There is a significant difference.

Hypothesis Test for 20000-40000, Mileage Range:
Test Statistic: -3.57072082691369
P-value: 0.0030719050336595337
Reject the null hypothesis: There is a significant difference.

Hypothesis Test for 40000-60000, Mileage Range:
Test Statistic: 23.519061069151917
P-value: 4.949769473983689e-23
Reject the null hypothesis: There is a significant difference.

Hypothesis Test for 60000-80000, Mileage Range:
Test Statistic: 27.997216590836068
P-value: 2.7123810201508093e-16
Reject the null hypothesis: There is a significant difference.

Hypothesis Test for 80000-100000, Mileage Range:
Test Statistic: 20.814680605018907
P-value: 4.739958812641532e-06
Reject the null hypothesis: There is a significant difference.

Hypothesis Test for 100000-120000, Mileage Range:
Test Statistic: nan
P-value: nan
Fail to reject the null hypothesis: There is no significant difference.

Hypothesis Test for 120000-140000, Mileage Range:
Test Statistic: nan
P-value: nan
Fail to reject the null hypothesis: There is no significant difference.

```

C:\Users\AAPAT\anaconda3\Lib\site-packages\numpy\core\fromnumeric.py:3464: RuntimeWarning: Mean of empty slice.
  return _methods._mean(a, axis=axis, dtype=dtype,
C:\Users\AAPAT\anaconda3\Lib\site-packages\numpy\core\_methods.py:192: RuntimeWarning: invalid value encountered in scalar divide
  ret = ret.dtype.type(ret / rcount)
C:\Users\AAPAT\anaconda3\Lib\site-packages\scipy\stats\_stats_py.py:1103: RuntimeWarning: divide by zero encountered in divide
  var *= np.divide(n, n-ddof) # to avoid error on division by zero
C:\Users\AAPAT\anaconda3\Lib\site-packages\scipy\stats\_stats_py.py:1103: RuntimeWarning: invalid value encountered in scalar multiply
  var *= np.divide(n, n-ddof) # to avoid error on division by zero

```

Q6.

```

In [44]: import pandas as pd

# Assuming df is your DataFrame with the relevant data
correlation_matrix = df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power', 'Price']].corr()

# Print the correlation matrix
print("Correlation Matrix:")
print(correlation_matrix)

```

Correlation Matrix:

	Distance driven (in miles)	Manufacturing year	\
Distance driven (in miles)	1.000000	-0.699032	
Manufacturing year	-0.699032	1.000000	
Engine_Power	-0.420251	0.601013	
Price	-0.765503	0.861719	

	Engine_Power	Price
Distance driven (in miles)	-0.420251	-0.765503
Manufacturing year	0.601013	0.861719
Engine_Power	1.000000	0.651781
Price	0.651781	1.000000

```


In [45]: sns.heatmap(correlation_matrix)

```

```

Out[45]: <Axes: >

```



Q7

```
In [47]: import statsmodels.api as sm
```



```
X= df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power']]
X= pd.get_dummies(X, drop_first= True)
y=df['Price']
X= sm.add_constant(X)
model= sm.OLS(y,X).fit()
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          Price    R-squared:                0.823
Model:                  OLS      Adj. R-squared:            0.817
Method:                 Least Squares    F-statistic:          148.4
Date:                  Thu, 14 Dec 2023    Prob (F-statistic):      6.30e-36
Time:                  02:17:33    Log-Likelihood:         -957.66
No. Observations:      100    AIC:                  1923.
Df Residuals:          96    BIC:                  1934.
Df Model:               3
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
-----	-----	-----	-----	-----	-----	-----
const	-6.251e+06	8.35e+05	-7.490	0.000	-7.91e+06	-4.59e+06
Distance driven (in miles)	-0.1002	0.019	-5.307	0.000	-0.138	-0.063
Manufacturing year	3105.1656	413.207	7.515	0.000	2284.957	3925.374
Engine_Power	28.7089	7.369	3.896	0.000	14.082	43.336
-----	-----	-----	-----	-----	-----	-----
Omnibus:	12.954	Durbin-Watson:	1.631			
Prob(Omnibus):	0.002	Jarque-Bera (JB):	20.953			
Skew:	0.547	Prob(JB):	2.82e-05			
Kurtosis:	4.958	Cond. No.	1.24e+08			
-----	-----	-----	-----	-----	-----	-----

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.24e+08. This might indicate that there are strong multicollinearity or other numerical problems.

Q7 Model Performance

```
In [49]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
```

```

X= df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power' ]]
X= pd.get_dummies(X, drop_first= True)
y= df['Price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

model = LinearRegression()

# Train the model on the training set
model.fit(X_train, y_train)

# Predict on the testing set
y_pred = model.predict(X_test)

# Evaluate model performance
mse= mean_squared_error(y_test, y_pred)
print(f'Mean Squared Error: {mse}')

coefficients = pd.DataFrame({'Variable': X.columns, 'Coefficient': model.coef_})
print(coefficients)

```

Mean Squared Error: 10284314.756285237

	Variable	Coefficient
0	Distance driven (in miles)	-0.105118
1	Manufacturing year	3080.958551
2	Engine_Power	28.452703

Q8.

```

In [57]: import statsmodels.api as sm
import matplotlib.pyplot as plt
import seaborn as sns

model.fit(X,y)

y_pred= model.predict(X)
residuals= y - y_pred

plt.figure(figsize= (8,4))
plt.subplot(1,2,1)

```


```
sns.scatterplot(x=y_pred, y=residuals)
plt.title('Residuals vs Predicted Values')
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.axhline(y=0, color='r', linestyle='--')
plt.show()

#Check for multicollinearity using VIF


from statsmodels.stats.outliers_influence import variance_inflation_factor

# Assuming X is Dataframe
vif_data= pd.DataFrame()
vif_data['Variable']= X.columns
vif_data['VIF']= [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]
print(vif_data)

plt.figure(figsize=(10, 6))
sns.histplot(residuals, kde=True)
plt.title('Distribution of Residuals')
plt.xlabel('Residuals')
plt.ylabel('Frequency')
plt.tight_layout()
plt.show()
```



	Variable	VIF
0	Distance driven (in miles)	4.837870
1	Manufacturing year	30.209231
2	Engine_Power	19.895675



```
In [59]: import statsmodels.api as sm
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd
from statsmodels.stats.outliers_influence import variance_inflation_factor

X= df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power']]
X= sm.add_constant(X)
```

```
y=df['Price']

model= sm.OLS(y, X).fit()

residuals= model.resid

#Linearity
plt.figure(figsize= (8,4))
plt.scatter(model.fittedvalues, residuals)
plt.title('Residuals vs Fitted Values (Linearity Check)')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.show()

#Independence
plt.figure(figsize= (8,4))
plt.plot(residuals)
plt.title('Residuals over time (Independence Check)')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.show()


#Homoscedasticity

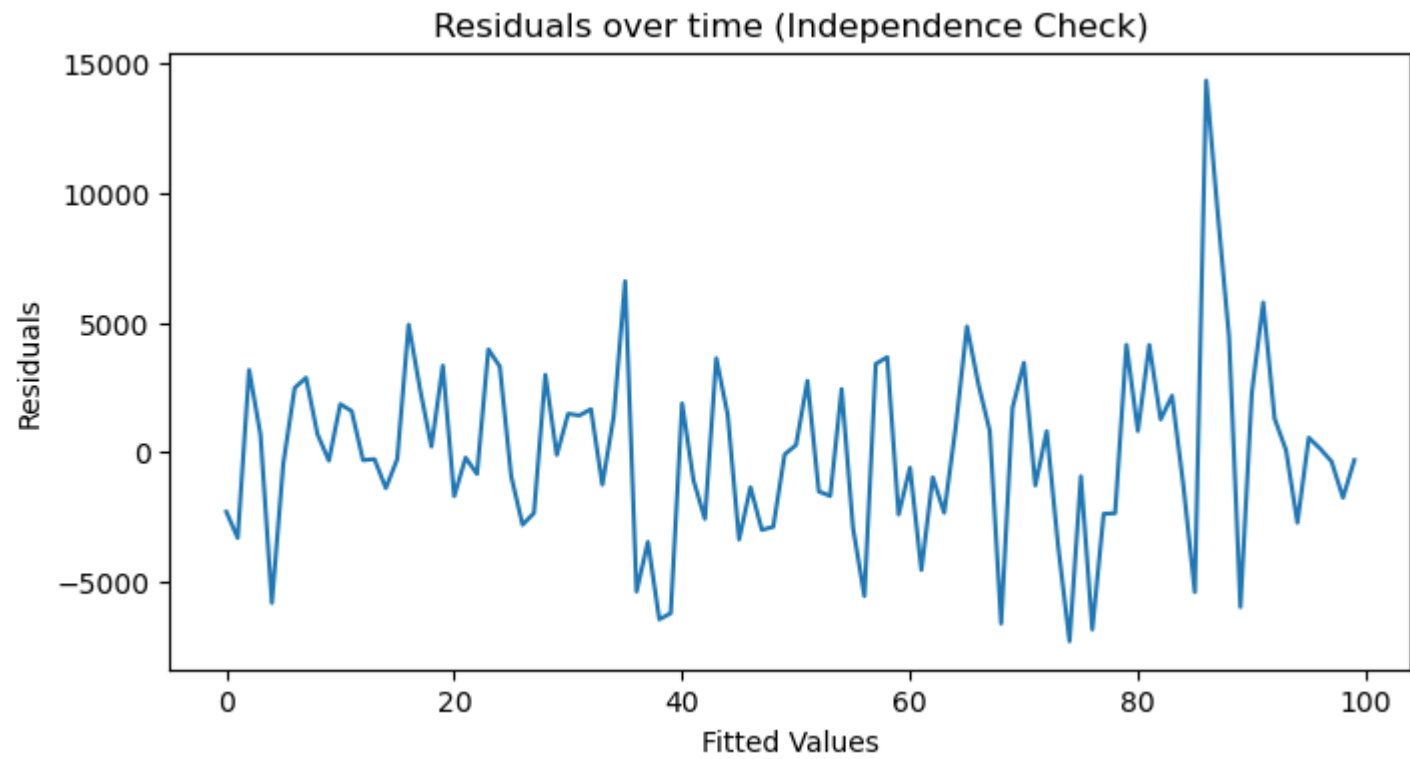
plt.figure(figsize=(8,4))
sns.scatterplot(x=model.fittedvalues, y=residuals)
plt.axhline(y=0, color='red', linestyle='--')
plt.title('Residuals vs Fitted Values (Homoscedasticity Check)')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.show()


#Normality of Residuals
plt.figure(figsize=(8,4))
sns.histplot(residuals, kde= True)
plt.title('Distribution of Residuals (Normality Check)')
plt.xlabel('Residuals')
plt.show()


#Multicollinearity
vif_data= pd.DataFrame()
```

```
vif_data['Variable'] = X.columns  
vif_data['VIF'] = [variance_inflation_factor(X.values, i) for i in range(X.shape[1])]  
  
print(vif_data)
```









	Variable	VIF
0	const	5.489399e+06
1	Distance driven (in miles)	1.955591e+00
2	Manufacturing year	2.520749e+00
3	Engine_Power	1.565477e+00

```
In [65]: import statsmodels.api as sm
import seaborn as sns
import matplotlib.pyplot as plt
import pandas as pd

# Assuming df is your DataFrame
# Replace column names with the actual column names from your dataset
# Select relevant independent variables

X= df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power']]

# Add a constant term to the independent variables/
```

```
X= sm.add_constant(X)

# Target variable

y = df['Price']

# Fit the Linear regression model


model = sm.OLS(y, X).fit()

# Get the residuals

residuals = model.resid

# Create a scatter plot of predicted values against residuals with a regression Line

plt.figure(figsize=(12, 6))
sns.regplot(x=model.fittedvalues, y=residuals, lowess=True, line_kws={'color': 'red'})
plt.title('Regression Line Plot for Residual Analysis')
plt.xlabel('Fitted Values')
plt.ylabel('Residuals')
plt.show()
```



```
In [70]: import statsmodels.api as sm
import pandas as pd

# Assuming df is your DataFrame containing the dataset
# Replace column names with the actual column names from your dataset
# Select relevant independent variables
X = df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power']]

# Add a constant term to the independent variables
X = sm.add_constant(X)

# Target variable
```

```
y = df['Price']

# Fit the Linear regression model
model = sm.OLS(y, X).fit()

# Now, let's use the model for prediction
# Example input data for prediction
new_data = {'Distance driven (in miles)': 50000, 'Manufacturing year': 2016, 'Engine_Power': 200}

# Create a DataFrame for the new data with the same columns as X
new_data_df = pd.DataFrame([new_data], columns=X.columns)

# Make the prediction without adding a constant term
predicted_price = model.predict(new_data_df)
print(f'Predicted Price: ${predicted_price.iloc[0]:.2f}')
```

Predicted Price: \$nan

```
In [71]: import statsmodels.api as sm
# Reputie coum nomes with the actual column names from your dataset
# Select relevant independent variables

X= df[['Distance driven (in miles)', 'Manufacturing year', 'Engine_Power']]

# Add a constant term to the independent variables
X= sm.add_constant (X)

# Target variable
y = df['Price']

# Print the summary

print(model.summary())
```

OLS Regression Results

=====						
Dep. Variable:	Price	R-squared:		0.823		
Model:	OLS	Adj. R-squared:		0.817		
Method:	Least Squares	F-statistic:		148.4		
Date:	Thu, 14 Dec 2023	Prob (F-statistic):		6.30e-36		
Time:	03:47:43	Log-Likelihood:		-957.66		
No. Observations:	100	AIC:		1923.		
Df Residuals:	96	BIC:		1934.		
Df Model:	3					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-6.251e+06	8.35e+05	-7.490	0.000	-7.91e+06	-4.59e+06
Distance driven (in miles)	-0.1002	0.019	-5.307	0.000	-0.138	-0.063
Manufacturing year	3105.1656	413.207	7.515	0.000	2284.957	3925.374
Engine_Power	28.7089	7.369	3.896	0.000	14.082	43.336
=====						
Omnibus:	12.954	Durbin-Watson:		1.631		
Prob(Omnibus):	0.002	Jarque-Bera (JB):		20.953		
Skew:	0.547	Prob(JB):		2.82e-05		
Kurtosis:	4.958	Cond. No.		1.24e+08		
=====						

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.24e+08. This might indicate that there are strong multicollinearity or other numerical problems.